

VT16 – EL2450 – Assignment II

Alexandros Filotheou
871108-5590
alefil@kth.se

Roberto Sanchez-Rey
840616-9139
rosr@kth.se

1 Part One: Scheduling

1.1 Rate Monotonic scheduling

1.1.1 Question 1

Rate Monotonic is an scheduling method that assigns fixed priorities to tasks, proportional to its activation frequency. That means that for any given tasks J_a, J_b with periods $T_a < T_b$, J_a is assigned a higher priority than J_b .

1.1.2 Question 2

A set of periodic tasks $\{J_i\}$ is schedulable with Rate Monotonic scheduling if

$$U = \sum_i \frac{C_i}{T_i} \leq n(2^{1/n} - 1)$$

In the case where $T_1 = 20, T_2 = 29, T_3 = 35$ ms and $C_i = 6$ ms, $i = \{1, 2, 3\}$, $U = 0.678$ and $n(2^{1/n} - 1) = 0.78$. Hence tasks J_1, J_2, J_3 are schedulable with RM.

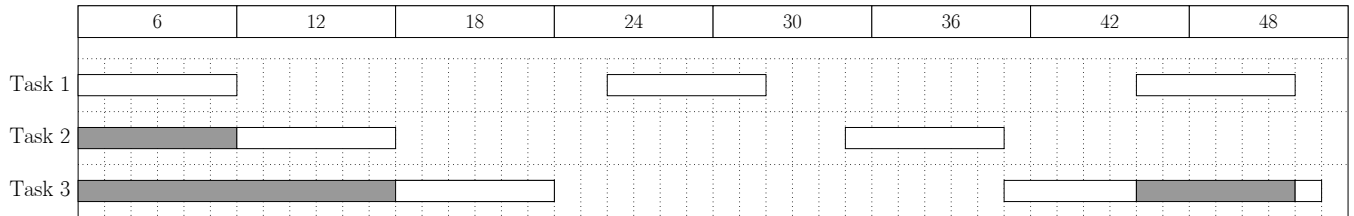


Figure 1: A portion of the RM schedule σ for tasks J_1, J_2, J_3 . Shaded areas denote the waiting time.

1.1.3 Question 3

All penduli are stable. We observe that the higher the natural frequency of a pentulum, the quicker the response is in its rise time, although with magnified overshoot. This makes sense since the higher the natural frequency of a pendulum, the lower its length and the more difficult it is

to stabilize, hence the control must be swift. Figure 2 shows the angular displacement of each pendulum.

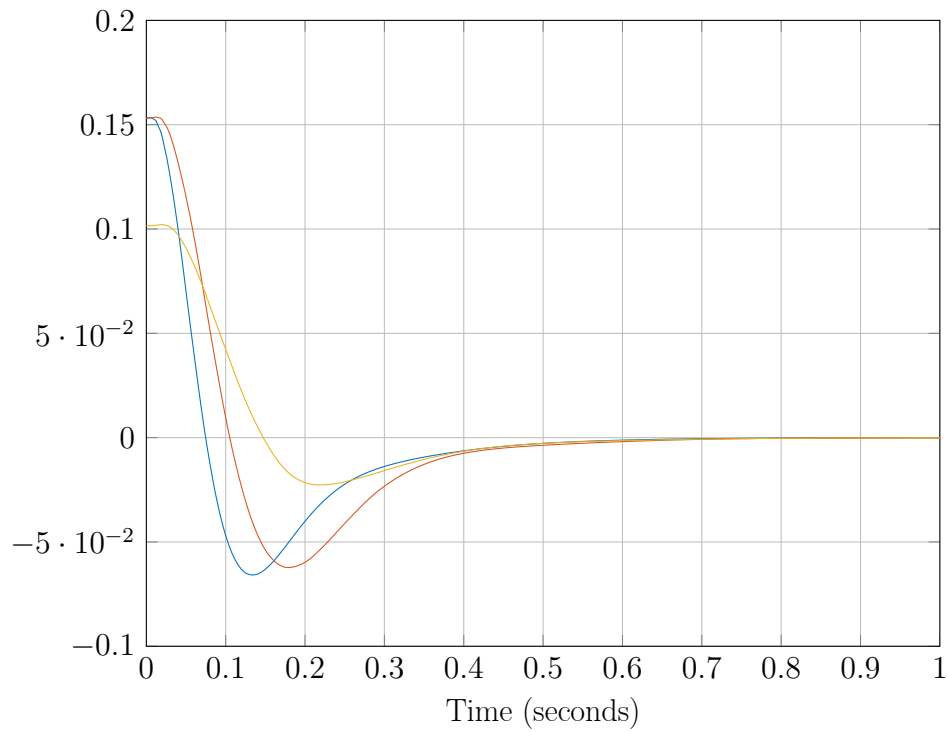


Figure 2: The angular displacement of each pendulum as a function of time. **Blue:** P_1 , **Red:** P_2 , **Orange:** P_3

1.1.4 Question 4

Figure 3 shows the actual execution schedule for each pendulum over a timespan of $lcm(20, 29, 35) = 4060$ ms: exactly one period of the schedule. As per the response to question 2, figure 4 illustrates that the schedule is indeed feasible by plotting the overall usage of the CPU over the aforementioned timespan.

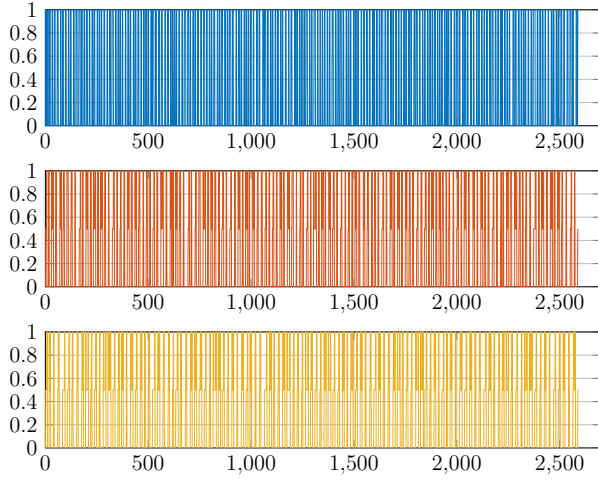


Figure 3: The calculated schedule for the three penduli. **Blue:** P_1 , **Red:** P_2 , **Orange:** P_3 . $C_i = 6$ ms.

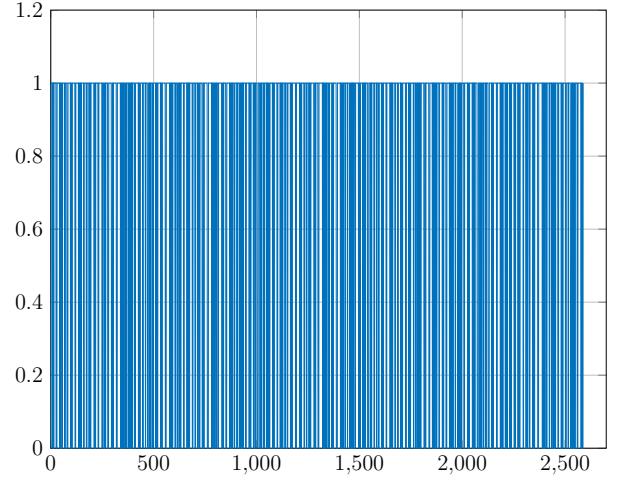


Figure 4: The overall processing usage. Notice that it is at most at 100%. $C_i = 6$ ms.

1.1.5 Question 5

In the case where $T_1 = 20, T_2 = 29, T_3 = 35$ ms and $C_i = 10$ ms, $i = \{1, 2, 3\}$, $U = 1.131 > 1$. Hence tasks J_1, J_2, J_3 are not schedulable under any scheduling scheme.

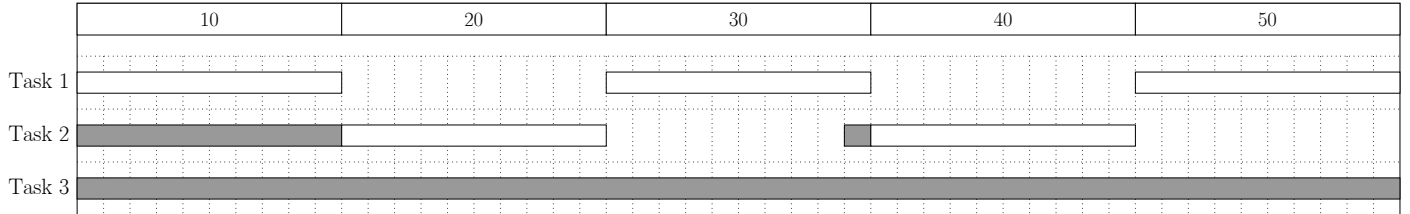


Figure 5: A portion of the RM schedule σ for tasks J_1, J_2, J_3 for $C_i = 10$ ms. Shaded areas denote the waiting time. Notice that J_3 misses its deadlines consecutively, indicative of the inability of schedulability.

Pendulum P_3 is left to its own devices and tends slowly but surely towards instability. This makes sense since no control signal pertaining to P_3 is assigned execution time in the processor. Figure 7 shows the angular displacement of pendulum P_3 as a function of time.

For penduli P_1 and P_2 , however, due to the increased execution time, delays are introduced, and the control input is not as swift as before, hence the increased magnitude of the overshoot and the larger rise and settling times. Figure 6 shows the angular displacement of each stable pendulum as a function of time.

Figures 8 and 9 show the angular displacement of penduli P_1 and P_2 respectively, as a function of time, for the two different cases of execution time.

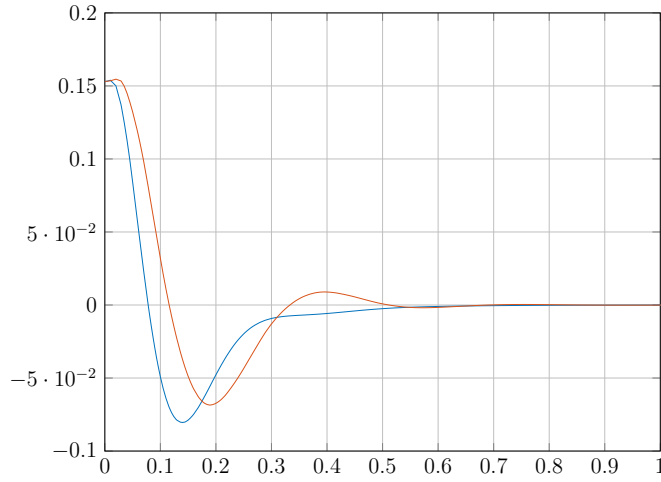


Figure 6: The angular displacement of the stable penduli P_1 and P_2 as a function of time. **Blue:** P_1 , **Red:** P_2 . $C_i = 10$ ms.

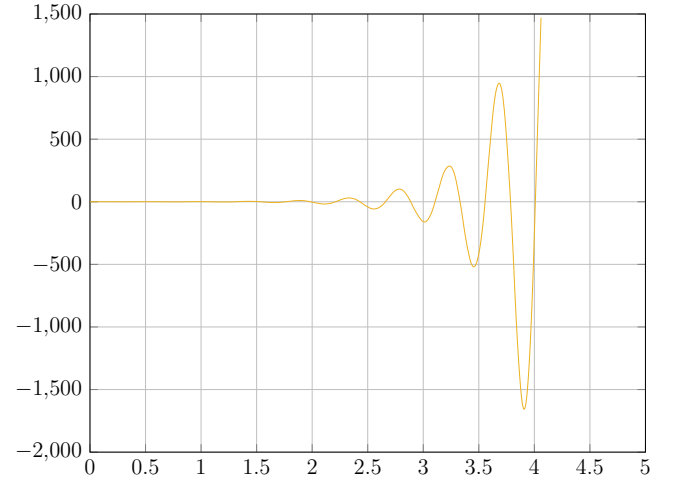


Figure 7: The angular displacement of the unstable pendulum P_3 as a function of time. $C_i = 10$ ms.

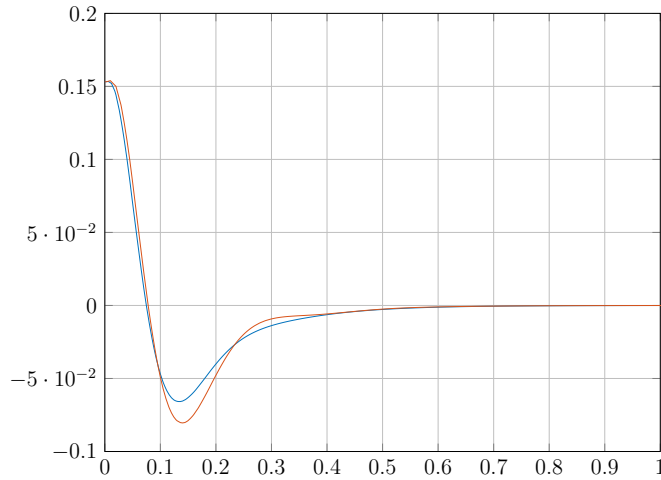


Figure 8: The angular displacement of pendulum P_1 as a function of time. **Blue:** $C_1 = 6$ ms, **Red:** $C_1 = 10$ ms

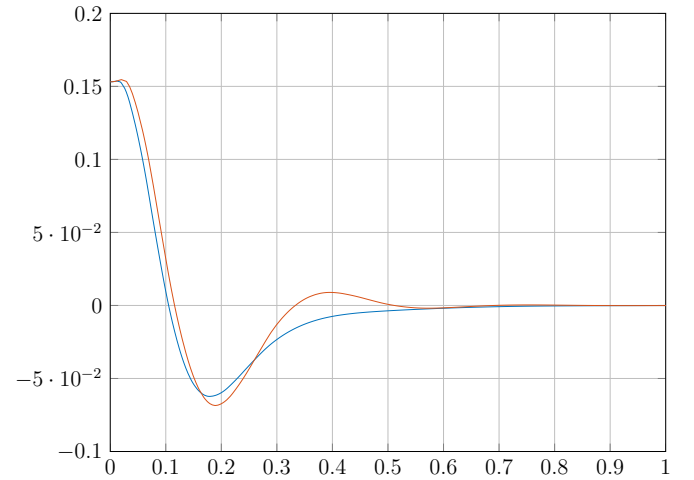


Figure 9: The angular displacement of pendulum P_2 as a function of time. **Blue:** $C_2 = 6$ ms, **Red:** $C_2 = 10$ ms

Figure 10 shows the schedule calculated for each pendulum with all jobs having execution time $C_i = 10$ ms. Figure 11 illustrates that the schedule is not feasible by plotting the overall usage of the CPU over the length of a schedule period, which is at all times 100%, indicative of the excessive processing load demanded.

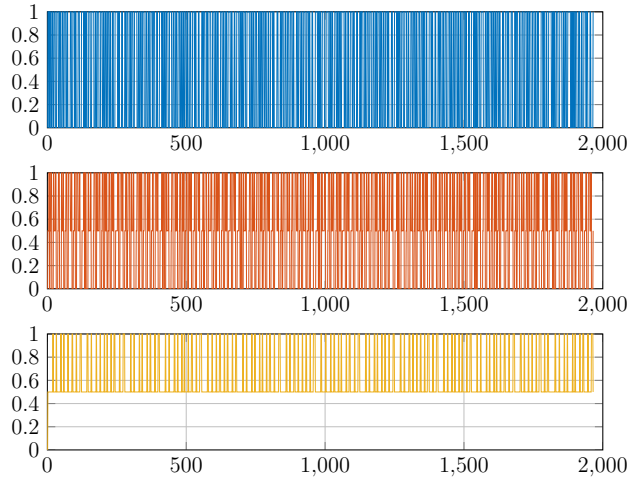


Figure 10: The calculated schedule for the three penduli. **Blue:** P_1 , **Red:** P_2 , **Orange:** P_3 . $C_i = 10$ ms.

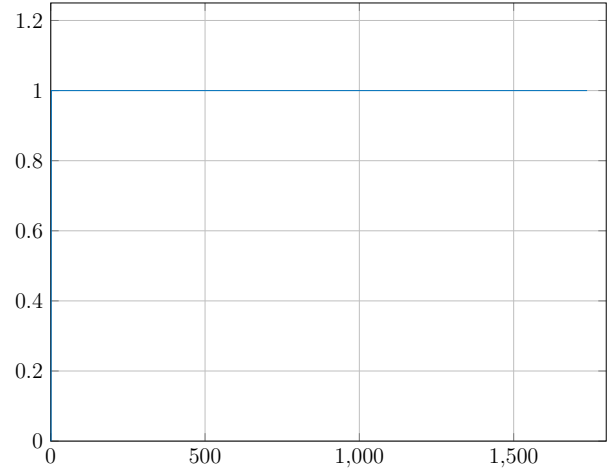


Figure 11: The overall processing usage. Notice that it is always at 100%. $C_i = 10$ ms.

1.2 Earliest Deadline First scheduling

1.2.1 Question 1

In contrast to Rate Monitoring scheduling, Earliest Deadline First scheduling assigns dynamic priorities to tasks, introducing a degree of flexibility. The task whose deadline is closest to the current timestep is given the highest priority and is executed for one time unit.

EDF is more lenient than RM, and this can be seen in the condition that identifies it:

$$U \leq 1 \Leftrightarrow \sigma \text{ is feasible}$$

whereas in RM

$$U \leq n(2^{1/n} - 1) (< 1 \text{ for } n \geq 2) \Rightarrow \sigma \text{ is feasible}$$

This means that if $U \leq 1$ for a certain collection of tasks $\{J_i\}$, there is always a feasible schedule for $\{J_i\}$, where the processor is utilized to the fullest it can be, whereas the former certainty does not hold with RM. However, if certain tasks are indeed more important than others and there are extra requirements per their execution, it is possible that EDF introduces delays between their release and start times due to the indiscrimination it shows to absolute task priorities.

1.2.2 Question 2

As stated above, a set of periodic tasks $\{J_i\}$ is schedulable under Earliest Deadline First scheduling if and only if

$$U = \sum_i \frac{C_i}{T_i} \leq 1$$

In the case where $T_1 = 20, T_2 = 29, T_3 = 35$ ms and $C_i = 6$ ms, $i = \{1, 2, 3\}$, $U = 0.678 \leq 1$. Hence tasks J_1, J_2, J_3 are schedulable with EDF.

1.2.3 Question 3

All penduli are stable for execution time $C_i = 6$ ms. We observe that the higher the natural frequency of a pendulum, the quicker the response is in its rise time, although with magnified overshoot. This makes sense since the higher the natural frequency of a pendulum, the lower its length and the more difficult it is to stabilize, hence the control must be swift. Figure 12 shows the angular displacement of each pendulum.

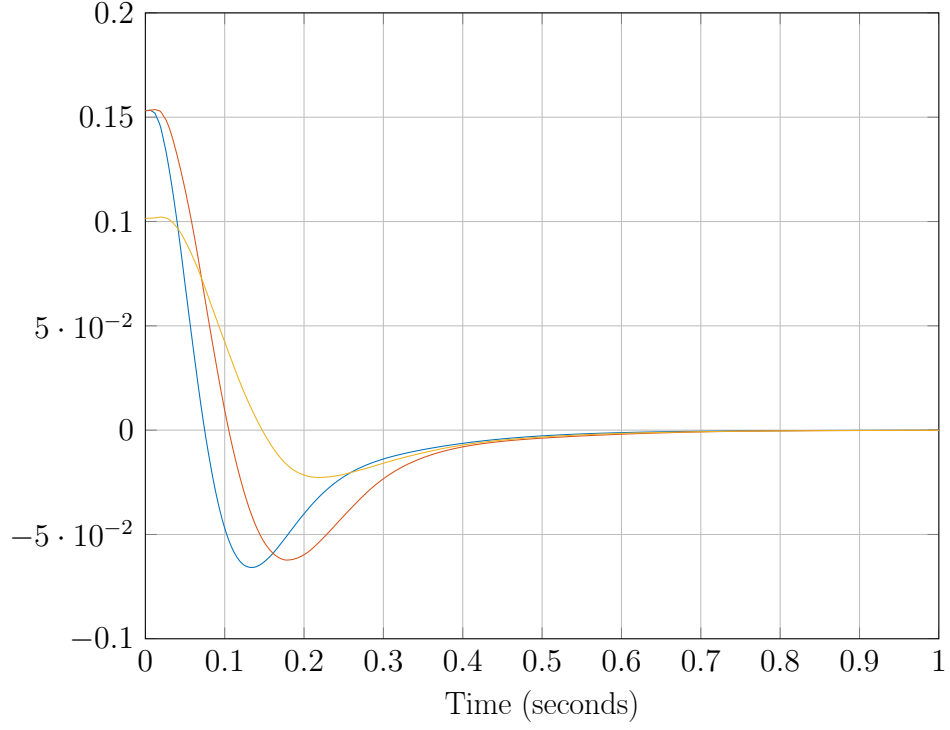


Figure 12: The angular displacement of each pendulum as a function of time. **Blue:** P_1 , **Red:** P_2 , **Orange:** P_3

1.2.4 Question 4

Figure 13 shows the schedule calculated for each pendulum over a timespan of $lcm(20, 29, 35) = 4060$ ms: exactly one period of the schedule. As per the response to question 2, figure 14 illustrates that the schedule is indeed feasible by plotting the overall usage of the CPU over the aforementioned timespan.

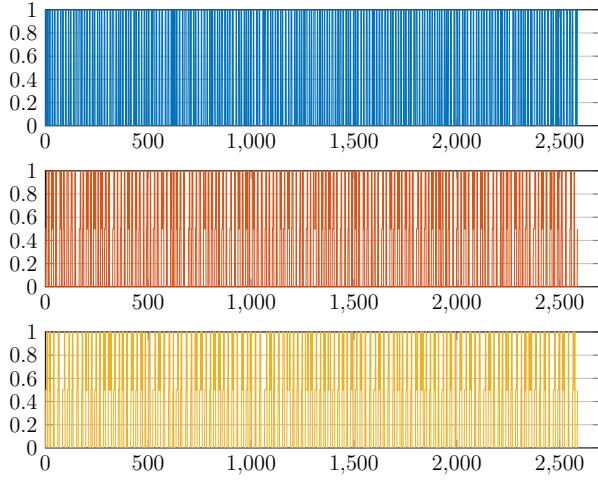


Figure 13: The calculated schedule for the three penduli. **Blue:** P_1 , **Red:** P_2 , **Orange:** P_3 . $C_i = 6$ ms.

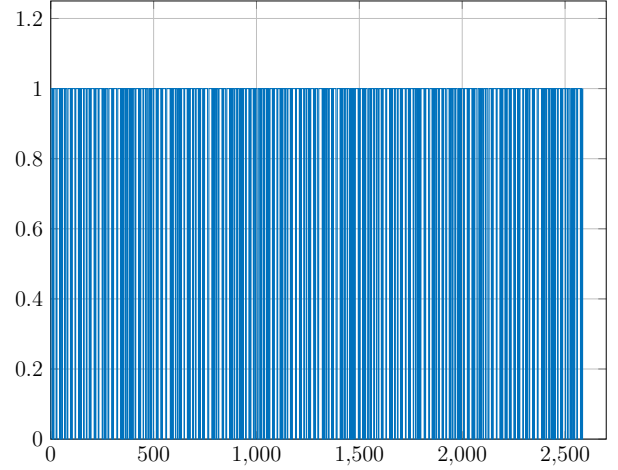


Figure 14: The overall processing usage. Notice that it is at most at 100%. $C_i = 6$ ms.

1.2.5 Question 5

In the case where $T_1 = 20, T_2 = 29, T_3 = 35$ ms and $C_i = 10$ ms, $i = \{1, 2, 3\}$, $U = 1.131 > 1$. Hence tasks J_1, J_2, J_3 are not schedulable under any scheduling scheme.

Here we note that the lack of feasibility of a schedule is connected to the notion and fact of deadlines not being met. This, however, does not mean that there can be no control over a system. All penduli are stable as opposed to the case of RM scheduling with the same execution time. This is because all control tasks can be preempted, hence every one of them is allowed execution in the processor, which results to all of them being executed. Hence, despite the fact that deadlines are missed, it is possible to control all three penduli, although it may be with worse control performance than in the case where $C_i = 6$ ms.

Due to the missing of deadlines, and, ultimately, due to the increased execution time, the control input is not as swift as before, hence the increased magnitude of the overshoot and the larger rise and settling times. Figure 15 shows the angular displacement of each pendulum as a function of time. Figures 16, 17 and 18 show the angular displacement of penduli P_1, P_2 and P_3 respectively, as a function of time, for the two different cases of execution time.

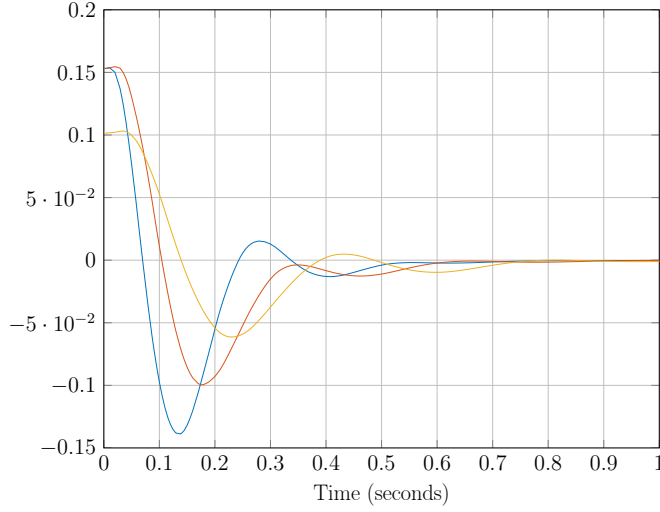


Figure 15: The angular displacement of each pendulum as a function of time. **Blue:** P_1 , **Red:** P_2 , **Orange:** P_3 . $C_i = 10$ ms.

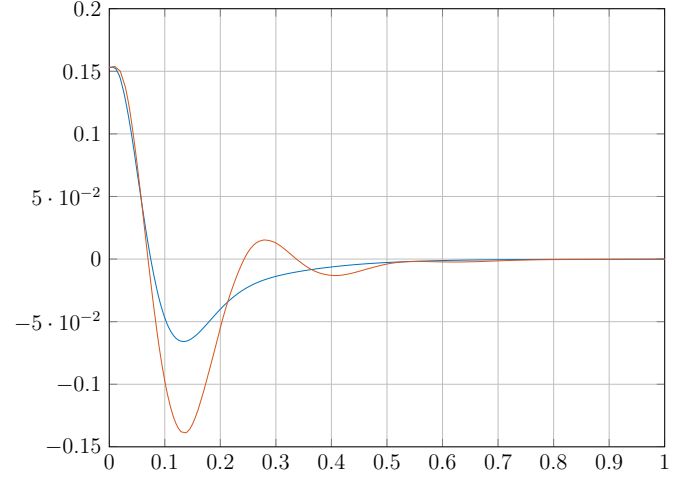


Figure 16: The angular displacement of pendulum P_1 as a function of time. **Blue:** $C_1 = 6$ ms, **Red:** $C_1 = 10$ ms

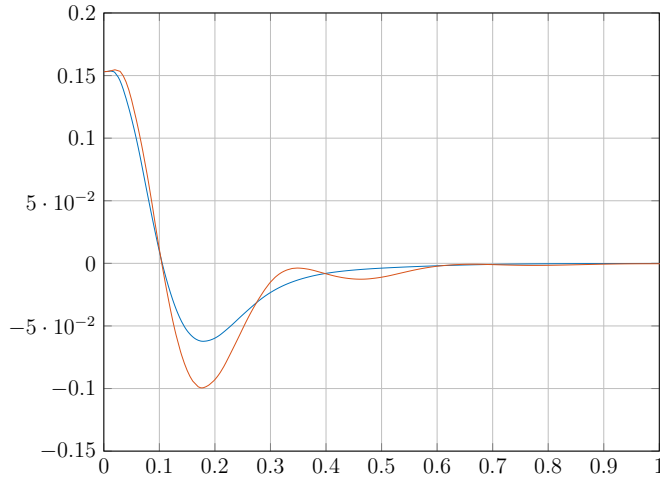


Figure 17: The angular displacement of pendulum P_2 as a function of time. **Blue:** $C_2 = 6$ ms, **Red:** $C_2 = 10$ ms

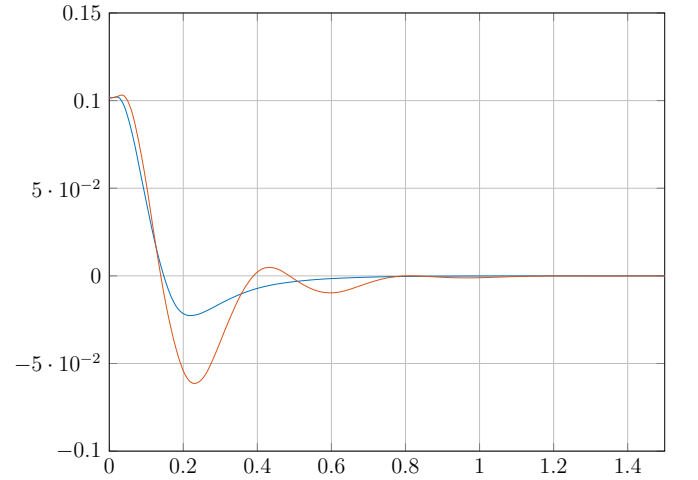


Figure 18: The angular displacement of pendulum P_3 as a function of time. **Blue:** $C_3 = 6$ ms, **Red:** $C_3 = 10$ ms

Figure 19 shows the schedule calculated for each pendulum with all tasks having execution time $C_i = 10$ ms. Figure 20 illustrates that the schedule is not feasible by plotting the overall usage of the CPU over the length of a schedule period, which is at all times 100%, indicative of the excessive processing load demanded.

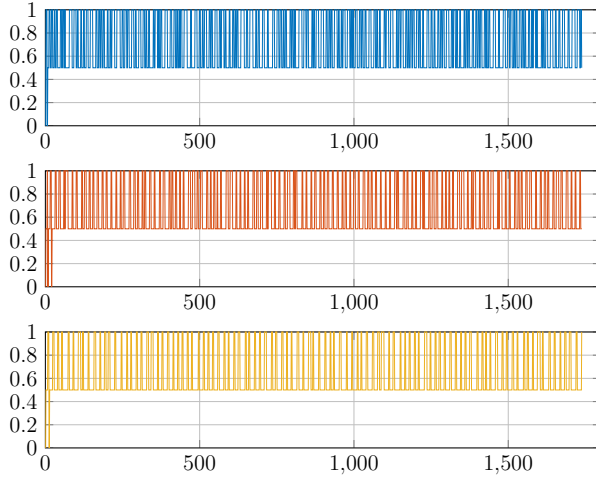


Figure 19: The calculated schedule for the three penduli. **Blue:** P_1 , **Red:** P_2 , **Orange:** P_3 . $C_i = 10$ ms.

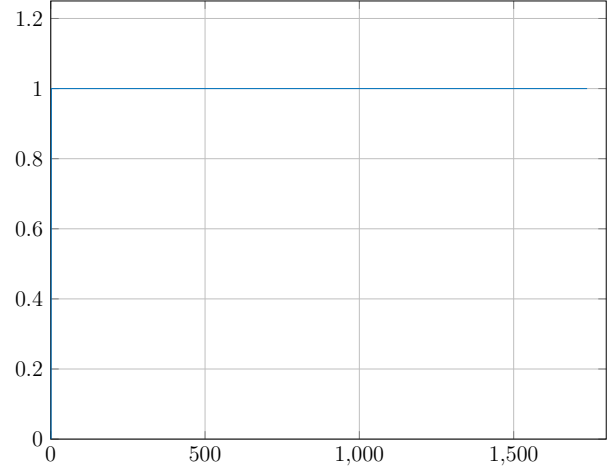


Figure 20: The overall processing usage. Notice that it is always at 100%. $C_i = 10$ ms.

1.2.6 Question 6

Figures 21, 23, 25, 27, and 29 feature the comparison of the course of the angular displacement of every stable pendulum, between the two considered scheduling policies. Figures 22, 24, 26, 28 and 30 graphically illustrate the evolution of the difference in angular displacement between the two considered scheduling policies for the two considered execution times for all stable penduli.

With execution time $C_i = 6$ ms the behaviour of all three penduli is nearly identical with respect to the two different scheduling policies. Scheduling under EDF makes the system have a higher overshoot but its peak difference from that of RM's is by an order of magnitude of 10^{-3} degrees. Hence their difference is negligible.

With execution time $C_i = 10$ ms there are two things to consider. First that pendulum P_3 is not controllable under RM but it is under EDF. Hence, if the purpose at hand is to control all three penduli, EDF is conclusively the only choice there is. On the other hand, though, EDF makes the system behave oscillatory in a degree larger than RM, with higher overshoot, rise and settling times.

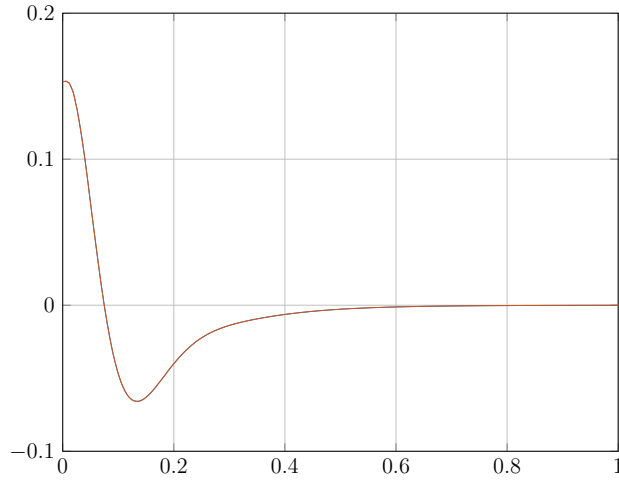


Figure 21: Evolution of the angular displacement of pendulum P_1 for $C_1 = 6$ ms. **Blue:** RM scheduling, **Red:** EDF scheduling

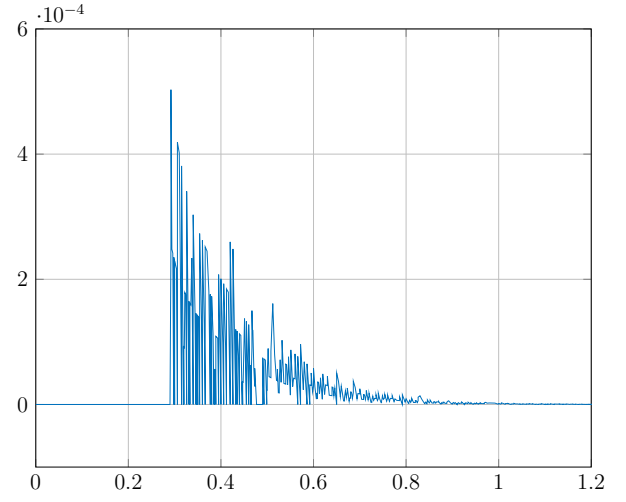


Figure 22: The evolution of the difference in angular displacement between RM and EDF of pendulum P_1 for execution time $C_1 = 6$ ms.

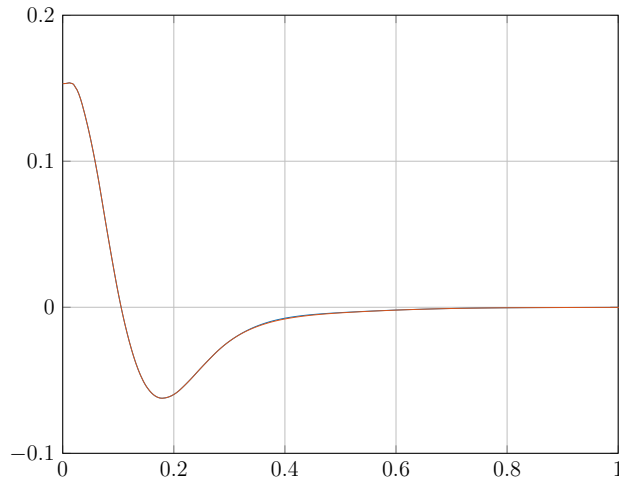


Figure 23: Evolution of the angular displacement of pendulum P_2 . $C_2 = 6$ ms. **Blue:** RM scheduling, **Red:** EDF scheduling

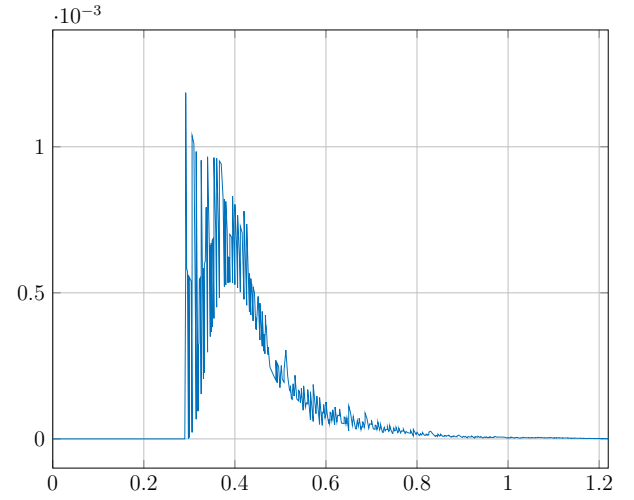


Figure 24: The evolution of the difference in angular displacement between RM and EDF of pendulum P_2 for execution time $C_2 = 6$ ms.

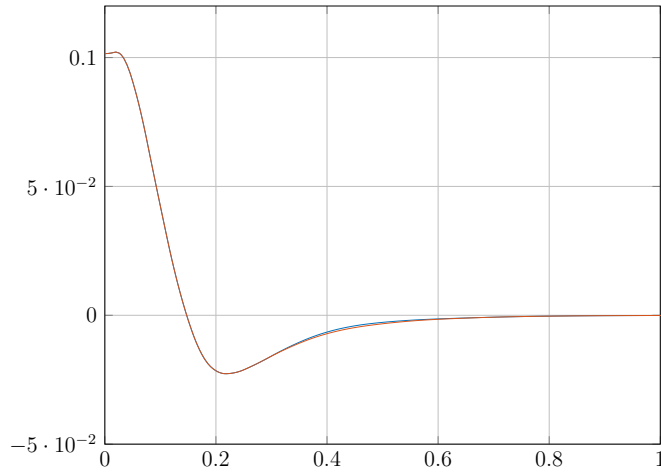


Figure 25: Evolution of the angular displacement of pendulum P_3 . $C_3 = 6$ ms. **Blue**: RM scheduling, **Red**: EDF scheduling

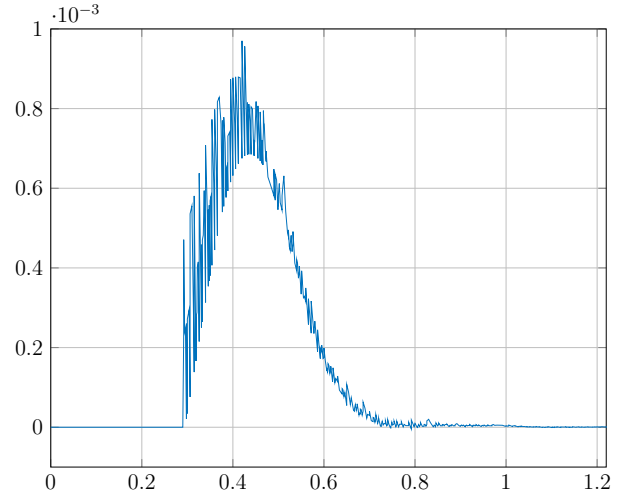


Figure 26: The evolution of the difference in angular displacement between RM and EDF of pendulum P_3 for execution time $C_3 = 6$ ms.

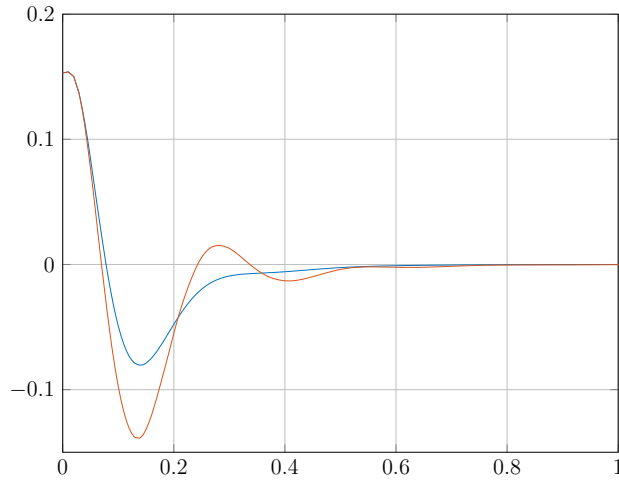


Figure 27: Evolution of the angular displacement of pendulum P_1 for $C_1 = 10$ ms. **Blue**: RM scheduling, **Red**: EDF scheduling

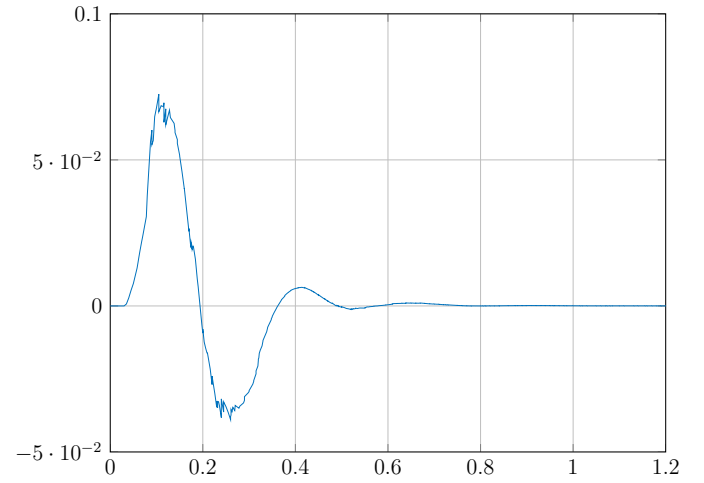


Figure 28: The evolution of the difference in angular displacement between RM and EDF of pendulum P_1 for execution time $C_1 = 10$ ms.

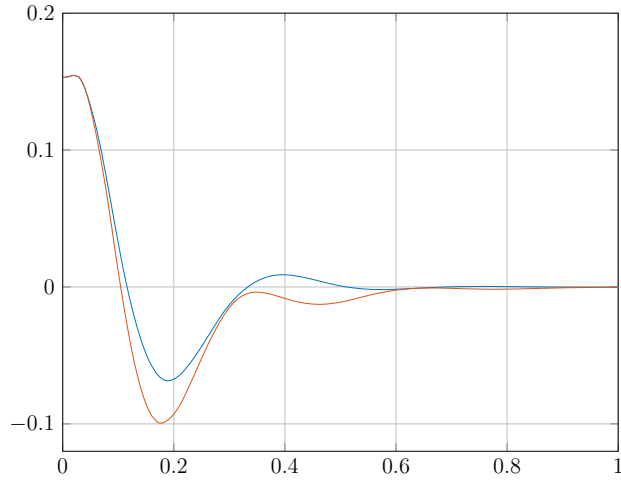


Figure 29: Evolution of the angular displacement of pendulum P_2 . $C_2 = 10$ ms. **Blue**: RM scheduling, **Red**: EDF scheduling

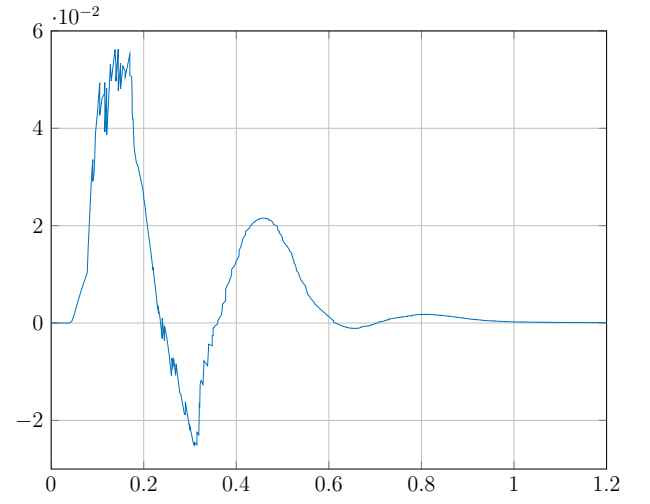


Figure 30: The evolution of the difference in angular displacement between RM and EDF of pendulum P_2 for execution time $C_2 = 10$ ms.